

Practical Machine Learning - Course Project

Data

First of all, the main data are read from the C drive and the first column with the number of observation is deleted.

```
RawData<-read.csv("C:/=Raw Data=/Practical Machine Learning Course Project/pml-training.csv",header=TRUE)
RawData<-RawData[,-1] #deleting number of observation (first column)
FinalTestData<-read.csv("C:/=Raw Data=/Practical Machine Learning Course Project/pml-testing.csv",header=TRUE)
FinalTestData<-FinalTestData[,-1] #deleting number of observation (first column)
n<-dim(RawData)[1]
```

The raw data set has dimensions 19622, 159. The raw data set should be incomplete, have missing values or not numerical values. The following code simply creates data set only with numerical values.

```
#numeric values
isNumber<-sapply(RawData,function(x) {sum(is.numeric(x) & !is.na(x))})
isNumber<-as.logical(isNumber==n)

#data with numeric values, displaying their means and names
CompleteData<-RawData[,isNumber]
outputVariable<-RawData$classes
apply(CompleteData,2,mean)
apply(CompleteData,2, var)
names(CompleteData)

#the same procedure with data for final testing
FinalTestCompleteData<-FinalTestData[,isNumber]
apply(FinalTestCompleteData,2,mean)
apply(FinalTestCompleteData,2, var)
names(FinalTestCompleteData)

sum(names(FinalTestCompleteData)!=names(CompleteData))

#output variable for CompleteData
y<-RawData$classes
```

The dimension of complete data set is 19622, 55. This is a significant decrease of attributes.

Basic Approach

The “fast and tidy” approach as a first step of the data analysis is usually recommended. This means try some fast, simple and straightforward way to analyze the data, evaluate it and find out how well/bad this approach is. Next, the improvements can be done based on those findings.

The PCA with knn classification and random forest were chosen as the “fast and tidy” approach on the complete data. The reasons are following.

- The number of attributes 55 is still high. PCA can reduce it.
- There is a high number of observations 19622. Even if we perform cross validation, the number will be still high. That is why knn-classification can perform well.
- Random forests are different kind of classifier, simple and effective.

Data Splitting

First of all, it is necessary to split the data into training set and testing set.

```
library(caret)
```

```
## Loading required package: lattice  
## Loading required package: ggplot2
```

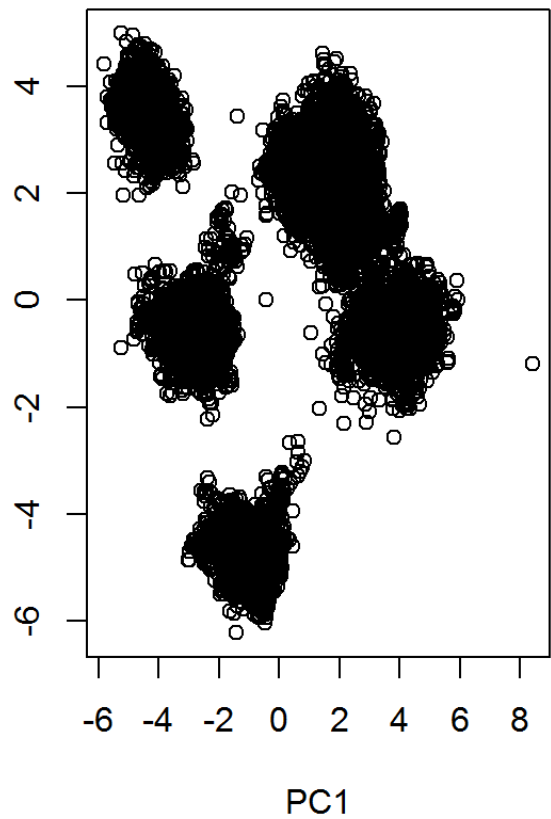
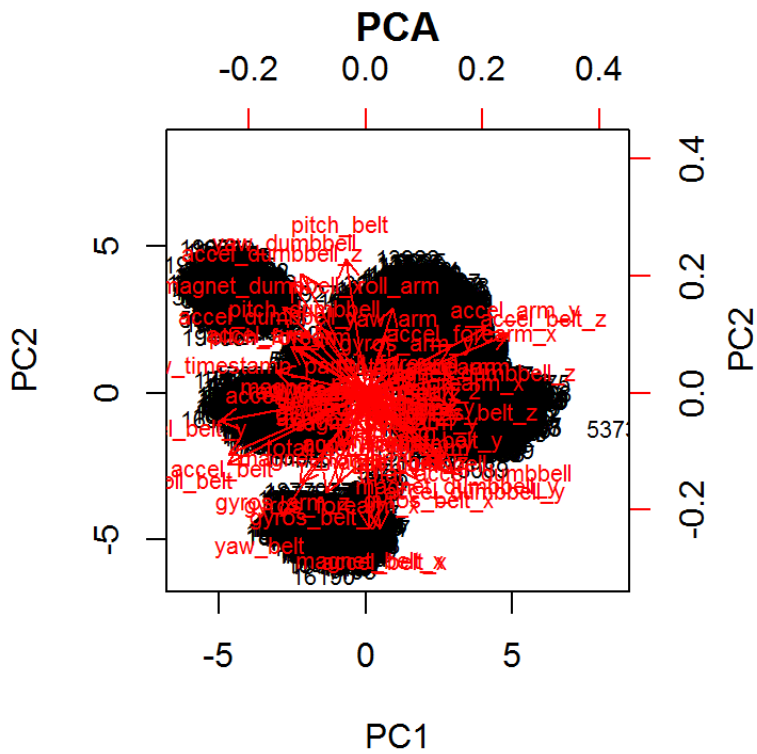
```
library(stats)  
set.seed(10201)  
inTrain<-createDataPartition(y=y,p=0.6,list=FALSE)  
#data splitting  
CompleteTrainingData<-CompleteData[inTrain,]  
CompleteTestingData<-CompleteData[-inTrain,]  
yTraining<-y[inTrain]  
yTesting<-y[-inTrain]  
#data standardisation  
CompleteTrainingData<-scale(CompleteTrainingData, center = TRUE, scale = TRUE)  
apply(CompleteTrainingData,2,mean)  
apply(CompleteTrainingData,2,sd)  
  
CompleteTestingData<-scale(CompleteTestingData, center = TRUE, scale = TRUE)  
  
FinalTestCompleteData<-scale(FinalTestCompleteData, center = TRUE, scale = TRUE)
```

PCA

Then, PCA is performed on the training set and biplot and screeplot are created.

```
PCA.CompleteTraining=prcomp(CompleteTrainingData, scale=FALSE)  
  
par(mfrow=c(1,2))  
biplot(PCA.CompleteTraining, scale=0,cex=.7,main="PCA") #zmÄ>na fontu  
plot(PCA.CompleteTraining$x[,1],PCA.CompleteTraining$x[,2],main="PCA.CompleteTraining",xlab="PC  
1",ylab="PC2")
```

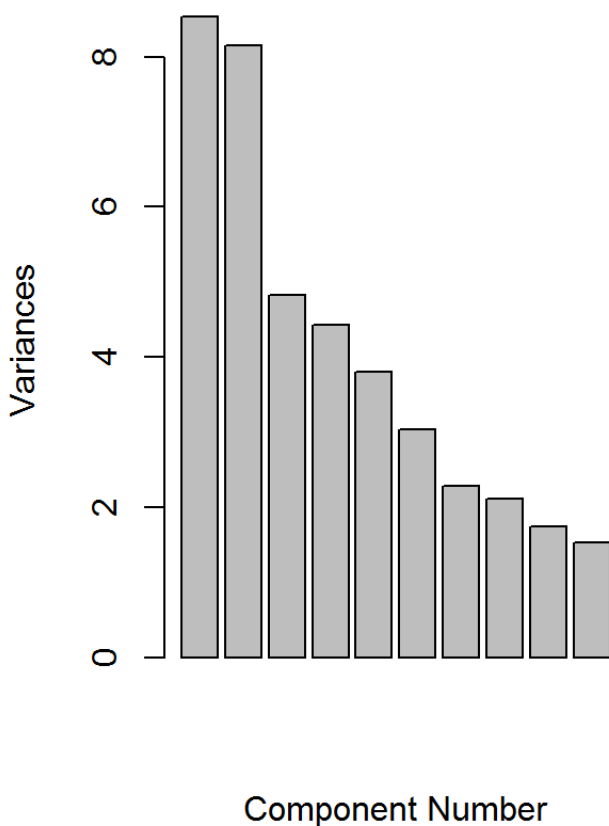
PCA.CompleteTraining



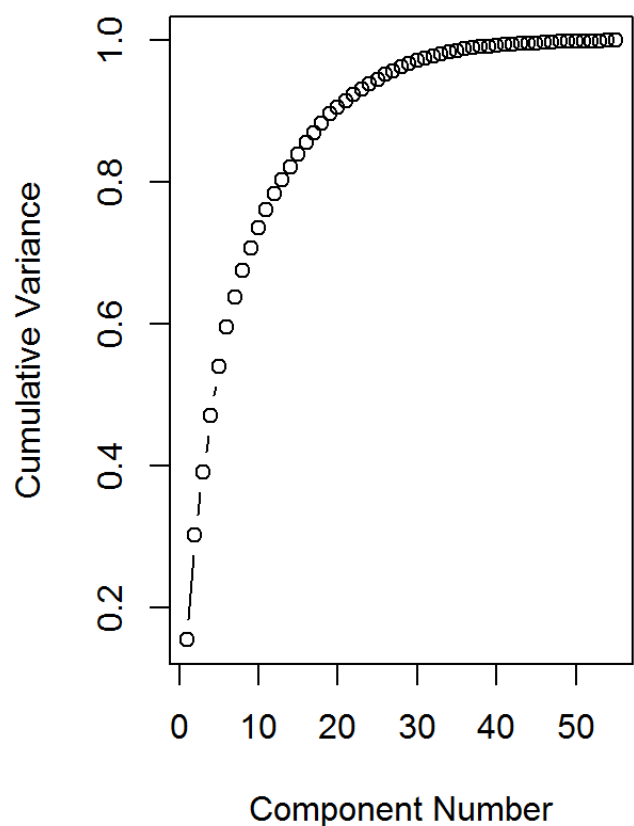
```
vars<-apply(PCA.CompleteTraining$x,2,var)
props<-vars/sum(vars)

par(mfrow=c(1,2))
plot(PCA.CompleteTraining,main="Screeplot",xlab="Component Number")
plot(cumsum(props),type="b",main="Cumulative Screeplot",xlab="Component Number",ylab="Cumulative Variance")
```

Screeplot



Cumulative Screeplot



```
par(mfrow=c(1,1))
```

K nearest neighbour

The K nearest neighbour classifier is performed on the data with two values of k, 5 and 7. The code for second one is hidden. For each value of k, the out-of-bag error on training data and classes for final data set are calculated.

```
#preparing data
nOfComp<-25
Dtrain<-PCA.CompleteTraining$x[,1:nOfComp]
Dtest<-CompleteTestingData%%PCA.CompleteTraining$rotation
Dtest<-Dtest[,1:nOfComp]

DFinalTest<-FinalTestCompleteData%%PCA.CompleteTraining$rotation
DFinalTest<-DFinalTest[,1:nOfComp]

#Learning
require(class)
```

```
## Loading required package: class
```

```
## Warning: package 'class' was built under R version 3.1.3
```

```

library(class)
k=5
PCA.knn<-knn(Dtrain,Dtest,yTraining,k=k)
#out-of-bag error - test data
KnnErr5<-1-sum(diag(table(PCA.knn,yTesting)))/sum(sum(table(PCA.knn,yTesting)))

#Final Test Prediction
FinalKnn5<-knn(Dtrain,DFinalTest,yTraining,k=k)

```

Random Forests

Finally, random forests are performed.

```

#preparing data
nOfComp<-25
Dtrain<-PCA.CompleteTraining$x[,1:nOfComp]
Dtest<-CompleteTestingData%%PCA.CompleteTraining$rotation
Dtest<-Dtest[,1:nOfComp]

DFinalTest<-FinalTestCompleteData%%PCA.CompleteTraining$rotation
DFinalTest<-DFinalTest[,1:nOfComp]

#learning
require(randomForest)

```

```

## Loading required package: randomForest
## randomForest 4.6-10
## Type rfNews() to see new features/changes/bug fixes.

```

```

PCA.randomForest=randomForest(x=Dtrain,y=yTraining)

#out-of-bag error - test data
randomForestErr<-1-sum(diag(table(predict(PCA.randomForest,Dtest),yTesting)))/sum(sum(table(predict(PCA.randomForest,Dtest),yTesting)))

#Final Test Prediction
FinalRandomForest<-predict(PCA.randomForest,DFinalTest)

```

Results

The out-of-bag error was 0.0453734 for knn with k=5, 0.0579913 for knn with k=7 and 0.0307163 for random forests. The predictions are following.

```

t(matrix(c(paste(FinalKnn5),paste(FinalKnn7),paste(FinalRandomForest)),ncol=3,nrow=20))

```

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13]
## [1,] "B"  "A"  "C"  "A"  "A"  "E"  "D"  "B"  "A"  "A"  "B"  "C"  "B"
## [2,] "B"  "A"  "A"  "A"  "A"  "E"  "D"  "D"  "A"  "A"  "B"  "C"  "B"
## [3,] "B"  "A"  "C"  "A"  "A"  "D"  "D"  "D"  "A"  "A"  "B"  "C"  "B"
##      [,14] [,15] [,16] [,17] [,18] [,19] [,20]
## [1,] "A"    "E"    "E"    "A"    "B"    "B"    "B"
## [2,] "A"    "E"    "E"    "A"    "B"    "B"    "B"
## [3,] "A"    "E"    "E"    "A"    "B"    "B"    "B"
```

The true values are B A B A A E D B A A B C B A E E A B B B.

The majority vote works well for 18 of 20 observation, it fails only for observations number 3 and 8. In this cases, there is needed some improvement.

This analysis shows, that even very simple approach can lead to reasonable results.